

# Table of Contents

---

- [Resources](#)
- [Hypothesis testing](#)
- [Type-I and Type-II errors and FP FN](#)
- [Normal distribution](#)
- [P-value](#)
- [H0 and H1](#)
- [Standard Error](#)
- [Margin of Error](#)
- [Pooled standard deviation and t-statistic](#)
- [F-test or ANOVA](#)
- [When to use Z-test and T-test](#)
- [Hypothesis Test Implementations](#)
  - [NOTES](#)
  - [One sample z-test](#)
  - [Coin Toss](#)
  - [z-test for coin toss](#)
  - [One sample t-test](#)
  - [Two sampled T-test](#)
  - [Paired sampled t-test](#)
  - [One-sampled Z-test](#)
  - [Two-sampled Z-test](#)
  - [F-test or ANOVA](#)
  - [Calculate zscore \(standardization\)](#)
  - [Binomial Test](#)
  - [Scipy normaltest \(D'Agostino and Pearson's test for normality\)](#)

## Resources

---

- [Youtube: 07e Python Data Analytics: Hypothesis Testing Interactive](#)
- [17 Hypothesis Testing in Python](#)
- [hypothesis-testing-spinning-the-wheel](#)
- [Analytics Vidhya: Everything you Should Know about p-value from Scratch for Data Science](#)
- [Hypothesis Tests](#)

## Hypothesis testing

---

Suppose a pizza place claims their delivery times are 30 minutes or less on average but you think it's more than that. So you conduct a hypothesis test and randomly sample some delivery times to test the claim:

- Null hypothesis — The mean delivery time is 30 minutes or less
- Alternative hypothesis — The mean delivery time is greater than 30 minutes

We'll use one-tailed test in our case since we only care about if the mean delivery time is greater than 30 minutes.

We use z-test for hypothesis testing:

$$z = \frac{\bar{x} - \mu}{\sigma / \sqrt{n}}$$

$$z = \frac{\hat{p} - p}{\sqrt{pq/n}} \quad (\hat{p} = h/n \text{ and } p = 0.5 \text{ for fair coin})$$

$z_{\text{crit}} = 1.96$  for two-tailed at  $\alpha = 0.05$   
 $= 1.645$  for one-tailed at  $\alpha = 0.05$

Remember right tail:

if  $z > z_{\text{crit}}$ : Reject  $H_0$  (our sample is too outside, another distribution will fit better.)

- suppose we think a coin is biased towards heads.
- We can conduct an experiment and toss the coin multiple times.
- Null hypothesis  $p=0.5$  (coin is fair) and Alternative hypothesis is coin is unfair (two-tailed test).
- Alternative hypothesis can also be that coin is biased towards tail ( $p>0.5$ ). This is called one-tailed test. (left-tailed or right-tailed test).
- We toss the coin multiple times (eg 10k times) and count number of heads.
- Calculate p-value and if p-value is less than 0.05, reject the null hypothesis.
- For example, if we get  $p\text{-value} = 0.375$  (two-sided), then it is not smaller than 0.05, so we fail to reject the null hypothesis.
- This means we do not have sufficient evidence to REJECT the belief of researcher coin is fair.
- This means, even if we get 4 heads out of 5 times, the coin is still unbiased at 95% confidence level.

The process of testing hypotheses can be compared to court trials. A person comes into court charged with a crime. A jury must decide whether the person is innocent (null hypothesis) or guilty (alternative hypothesis). Even though the person is charged with the crime, at the beginning of the trial (and until the jury declares otherwise) the accused is assumed to be innocent. Only if overwhelming evidence of the person's guilt can be shown is the jury expected to declare the person guilty--otherwise the person is considered innocent.

## Type-I and Type-II errors and FP FN

---

		Decision	
		Reject $H_0$	Accept $H_0$
State of the world	$H_0$ true	<b>A</b> Type I error	<b>B</b> Correct
	$H_0$ False	<b>C</b> Correct	<b>D</b> Type II error

In the jury trial there are two types of errors: (1) the person is innocent but the jury finds the person guilty, and (2) the person is guilty but the jury declares the person to be innocent. In our system of justice, the first error is considered more serious than the second error. These two errors along with the correct decisions are shown in the next table where the jury decision is shown in bold on the left margin and the true state of affairs is shown in bold along the top margin of the table.

	Truth is Person Innocent	Truth is Person Guilty
Jury Decides Person Innocent	Correct Decision	Type II Error
Jury Decides Person Guilty	Type I Error	Correct Decision

	In Fact $H_0$ is True	In Fact $H_0$ is False
Test Decides $H_0$ True	Correct Decision	Type II Error
Test Decides $H_0$ False	Type I Error	Correct Decision

In a jury trial the person accused of the crime is assumed innocent at the beginning of the trial, and unless the jury can find overwhelming evidence to the contrary, should be judged innocent at the end of the trial. Likewise, in hypothesis testing, the null hypothesis is assumed to be true, and unless the test shows overwhelming evidence that the null hypothesis is not true, the null hypothesis is accepted.

$H_0$ : Not producing more defectives

$H_1$ : Producing more defectives (Federal Checker wants to test it.)

Type I = False Positive

Reject  $H_0$ , when  $H_0$  is true.


Say producing more defects, when it is not.

Type II = False Negative

Fail to reject  $H_0$ , when  $H_1$  is true.


Say NOT producing more defects, when it is.

# Normal distribution





## P-value







The p-value consists of 3 parts: Is a coin biased if two out of two tosses are head? Answer: No. ( $p\text{-val} \leq \alpha$ )



p-value for **2 Heads** =  $0.25 + 0.25 + 0 = 0.5$

1st Flip    2nd Flip

Outcomes	Probability
 	0.25
 	0.5
 	0.25


We add **Part 2**, the probability of something else that is equally rare, because although getting **2 Heads** might seem special, it doesn't seem as special when we know that other things are just as rare.

**A p-value** is composed of three parts:

- 1) The probability random chance would result in the observation.
- 2) The probability of observing something else that is equally rare.
- 3) The probability of observing something rarer or more extreme.


7:26 / 25:14 • p-value defined as the sum of three parts >

Is a coin biased if two out of two tosses are head? Answer: No. ( $p\text{-val} \leq \alpha$ )



In other words, the data, getting **4 Heads** and **1 Tails**, did not convince us that our coin was special.

$$\frac{5}{32} + \frac{5}{32} + \frac{2}{32} = 0.375$$



The lower the p-value, the more surprising the evidence is, the more ridiculous our null hypothesis looks.

If the p-value is lower than a predetermined significance level then we reject the null hypothesis.

Now that we've collected some sampled delivery times, we perform the calculation and find that the mean delivery time is longer by 10 minutes with a p-value of 0.03.

What this means is that in a world where the pizza delivery time is 30 minutes or less (null hypothesis is true), there's a 3% chance we would see the mean delivery time is at least 10 minutes longer due to random noise.

The lower the p-value, the more meaningful the result because it is less likely to be caused by noise. There's a common misinterpretation of p-value for most people in our case: The p-value 0.03 means that there's 3% (probability in percentage) that the result is due to chance — which is not true.

P-value caveats:

The p-value only says if two groups A and B are different, it does not say how much different are they? Small p-value does not mean they are more different, if we change sample size even if % of effective drug in A, B might be similar, p-value can be small. This means in linear regression, we can NOT have feature importance based on p-values.

**Description:** p-value is the probability value that we select equal or more extreme results from the distribution. For example, if we toss the coin 5 times and got 4 heads, the p-value is  $P(4H1T) + P(4T1H) + P(5H) + P(5T) = 5/32 + 5/32 + 1/32 + 1/32 = 12/32 = 0.375$ . (note  $32 = 2^5$ ). The p-value tells how extreme is our result. We can also get pvalue using binomial test from scipy.

- Note: 4H1T means having one T in one of the five draws, and others are Heads, so chances are 5 (draw1, draw2, draw3, draw4, draw5).
- `scipy.stats.binomtest(k, n, p=0.5, alternative='two-sided')`

```

from scipy import stats
# binom_test is deprecated use binomtest
result = stats.binomtest(k=4, n=5, p=0.5, alternative='two-sided') # 0.375
pval = result.pvalue

print(result.proportion_estimate) # k/n
print(result.proportion_ci(confidence_level=0.95))

4H1T is 4 heads and 1 tail = 5C4 = 5!/(4!1!) = 5
HHHHT (place T in one of 5 places)
HHHHT
HHHTH
HHTHH
HTHHH
THHHH

```

Usually we compare the p-value with the user defined significance level alpha. If p-value is less than alpha, then we reject the null hypothesis and say that we have sufficient evidence to reject the null hypothesis in support of alternative hypothesis. In this case 0.375 is NOT less than 0.05 (95% confidence level) and we say that coin is fair. (we FAIL to reject null hypothesis) where null hypothesis is "coin is fair p=0.5".

## NOTES

When p-value is less than alpha, we reject the null hypothesis. In other words we have evidence to believe that the sample statistic is significantly different from the population parameter.

## WARNINGS

The p-value only says if two groups A and B are different, it does not say how much different are they? Small p-value does not mean they are more different, if we change sample size even if % of effective drug in A, B might be similar, p-value can be small. This means in linear regression, we can NOT have feature importance based on p-values.

# H0 and H1

---

H1 is the alternative hypothesis (also called research hypothesis)  
Usually, it is the hypothesis desired by the researcher.

e.g.

H1: The average score of class is better than national average.

H1: This particular coin is biased.

if  $p \leq \alpha$ : We reject the null hypothesis  
and say our research hypothesis is statistically significant.

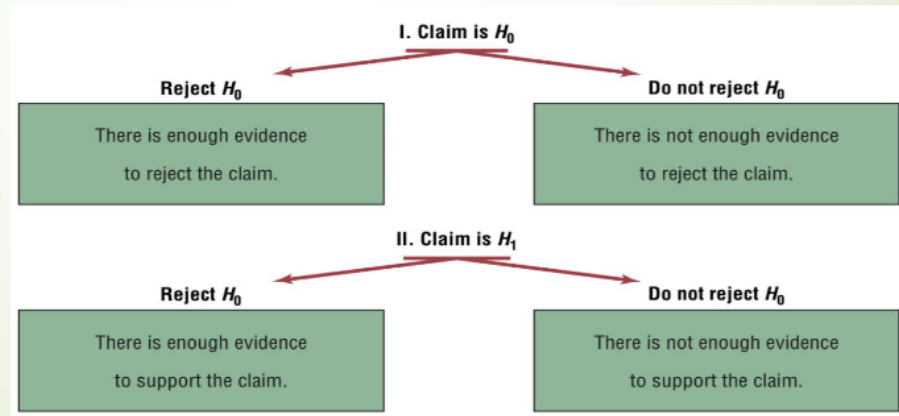
But, in `scipy` testing it is opposite:

H0: Two samples comes from normal distribution (Shapiro-Wald Test)

if  $p > \alpha$ : we fail to reject null hypothesis and say data is NORMAL.

## IMPORTANT NOTE:

- ▶ **When the null hypothesis is not rejected, we do not accept it as true. There is merely not enough evidence to say that it is false.**
- ▶ We conclude the alternative hypothesis (when we reject the null) because the data clearly support that conclusion.



## Standard Error

---

- Stat Trek: What is the Standard Error?

Table 1: Mean and Standard Deviation				
Measure Name	Symbol for Population	Symbol for Sample	Computation for Population	Computation for Sample
Mean	$\mu$	$\bar{x}$	$\mu = \frac{\sum_{i=1}^N x_i}{N}$	$\bar{x} = \frac{\sum_{i=1}^n x_i}{n}$
Standard Deviation	$\sigma$	$s_x$	$\sigma = \sqrt{\frac{\sum_{i=1}^N (x_i - \mu)^2}{N}}$ <p>or the equivalent form</p> $\sigma = \sqrt{\frac{\sum_{i=1}^N (x_i^2) - \frac{(\sum_{i=1}^N x_i)^2}{N}}{N}}$	$s_x = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n - 1}}$ <p>or the equivalent form</p> $s_x = \sqrt{\frac{\sum_{i=1}^n (x_i^2) - \frac{(\sum_{i=1}^n x_i)^2}{n}}{n - 1}}$

Standard Error of Sample Estimates

Sadly, the values of population parameters are often unknown, making it impossible to compute the standard deviation of a statistic. When this occurs, use the standard error.

The standard error is computed from known sample statistics. The table below shows how to compute the standard error for simple random samples, assuming the population size is at least 20 times larger than the sample size.

Statistic	Standard Error
Sample mean, $\bar{x}$	$SE_{\bar{x}} = s / \text{sqrt}(n)$
Sample proportion, $p$	$SE_p = \text{sqrt}[p(1 - p) / n]$
Difference between means, $\bar{x}_1 - \bar{x}_2$	$SE_{\bar{x}_1 - \bar{x}_2} = \text{sqrt}[s^2_1 / n_1 + s^2_2 / n_2]$
Difference between proportions, $\bar{p}_1 - \bar{p}_2$	$SE_{\bar{p}_1 - \bar{p}_2} = \text{sqrt}[p_1(1 - p_1) / n_1 + p_2(1 - p_2) / n_2]$

The equations for the standard error are identical to the equations for the standard deviation, except for one thing - the standard error equations use statistics where the standard deviation equations use parameters. Specifically, the standard error equations use  $p$  in place of  $P$ , and  $s$  in place of  $\sigma$ .



## Standard Deviation of Sample Estimates

Statisticians use sample statistics to estimate population [parameters](#). Naturally, the value of a statistic may vary from one sample to the next.

The variability of a statistic is measured by its standard deviation. The table below shows formulas for computing the standard deviation of statistics from [simple random samples](#). These formulas are valid when the population size is much larger (at least 20 times larger) than the sample size.

Statistic	Standard Deviation
Sample mean, $\bar{x}$	$\sigma_{\bar{x}} = \sigma / \sqrt{n}$
Sample proportion, $p$	$\sigma_p = \sqrt{P(1-P) / n}$
Difference between means, $\bar{x}_1 - \bar{x}_2$	$\sigma_{\bar{x}_1 - \bar{x}_2} = \sqrt{\sigma_1^2 / n_1 + \sigma_2^2 / n_2}$
Difference between proportions, $\bar{p}_1 - \bar{p}_2$	$\sigma_{\bar{p}_1 - \bar{p}_2} = \sqrt{P_1(1-P_1) / n_1 + P_2(1-P_2) / n_2}$

**Note:** In order to compute the standard deviation of a sample statistic, you must know the value of one or more population parameters. For example, to compute the standard deviation of the sample mean ( $\sigma_{\bar{x}}$ ), you need to know the variance of the population ( $\sigma$ ).

Multiplier Number ( $z^*$ )	Level of Confidence
3.0	99.7%
2.58 (2.576)	99%
2.0 (more precisely 1.96)	95%
1.645	90%
1.282	80%
1.15	75%
1.0	68%

## Margin of Error

---

## Formula

$$\text{MOE}_{\gamma} = z_{\gamma} \times \sqrt{\frac{\sigma^2}{n}}$$

MOE = margin of error

$\gamma$  = confidence level

$z_{\gamma}$  = quantile

$\sigma$  = standard deviation

$n$  = sample size

Note:  $z_{\gamma}$  is 1.96 for two tailed test with 95% confidence interval.

## Pooled standard deviation and t-statistic

---

$$s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}}$$

$$t = \frac{(\bar{x}_1 - \bar{x}_2) - (\mu_1 - \mu_2)}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

## F-test or ANOVA

---

When we have more than two groups we use ANOVA test. In ANOVA one column is categorical and target must be numerical (hence the name one-way anova). Categorical features are also called factor.

## When to use Z-test and T-test

---

Some notes:

- Use Z-test when your sample size is greater than 30. Otherwise, use a t test.

- Data points should be independent from each other. In other words, one data point isn't related or doesn't affect another data point.
- Your data should be normally distributed. However, for large sample sizes (over 30) this doesn't always matter.
- Your data should be randomly selected from a population, where each item has an equal chance of being selected.
- Sample sizes should be equal if at all possible.

## Hypothesis Test Implementations

---

### NOTES

- `scipy.stats` has `ttest` and `binom` test but it does not have `ztest`.
- We need to get `ztest` from `statsmodels.stats.weightstats` as `wstats` and then `wstats.ztest`

### One sample z-test

```
z = xbar - mu / (sigma/sqrt(n))

zstar = 1.645 for one-tailed test at alpha = 5% (gamma = 95%)
        1.96  for two-tailed test at alpha = 5% (gamma = 95%)

if z > zstar: Reject H0
if p < 0.05 : Reject H0
```


A principal at a certain school claims that the students in his school are above average intelligence. A random sample of thirty students IQ scores have a mean score of 112.5. Is there sufficient evidence to support the principal's claim? The mean population IQ is 100 with a standard deviation of 15.

**Step 1:** State the **Null hypothesis**. The accepted fact is that the **population mean** is 100, so:  $H_0: \mu=100$ .

**Step 2:** State the **Alternate Hypothesis**. The claim is that the students have above average IQ scores, so:  
 $H_1: \mu > 100$ .

The fact that we are looking for scores "greater than" a certain point means that this is a **one-tailed test**.

**Step 3:** Draw a picture to help you visualize the problem.



**Step 4:** State the **alpha level**. If you aren't given an **alpha level**, use 5% (0.05).

**Step 5:** Find the **rejection region** area (given by your alpha level above) from the **z-table**. An area of .05 is equal to a **z-score** of 1.645.

**Step 6:** Find the **test statistic** using this formula:

$$Z = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}$$

For this set of data:  $z = (112.5 - 100) / (15 / \sqrt{30}) = 4.56$ .

**Step 6:** If Step 6 is greater than Step 5, reject the null hypothesis. If it's less than Step 5, you cannot reject the null hypothesis. In this case, it is greater ( $4.56 > 1.645$ ), so you can reject the null.

## Coin Toss

### z-test for coin toss

- Ref: [ipython-books.github.io72-getting-started-with-statistical-hypothesis-testing-a-simple-z-test/](https://github.com/ipython-books/72-getting-started-with-statistical-hypothesis-testing-a-simple-z-test/)
- `scipy.stats` has t-test but it does not have ztest, to do ztest we need to use `statsmodels.stats` and then `weightstats`

- For coin toss, we can also use `scipy.stats`.

Let's suppose that after  $n=100$  flips, we get  $h=61$  heads.  
 We choose a significance level of  $0.05$ : is the coin fair or not?  
 Our null hypothesis is: the coin is fair ( $p=1/2$ ).  
 $H_0$  : a coin is a fair coin.  
 $H_1$  : a coin is a tricky coin.  
 $\alpha = 5\%$  or  $0.05$

```
import numpy as np
import scipy.stats as stats
import scipy.special as ssp
n = 100 # number of coin flips
h = 61 # number of heads
p = 0.5 # null-hypothesis of fair coin

# find z-statistic
xbar = float(h) / n

var = p*(1-p) # variance is npq but in the book, it is pq
sigma = np.sqrt(var)
expected_mean = p # mu = np in wikipedia, but p in book.
std_err = sigma/np.sqrt(n)
z = (xbar - expected_mean) / std_err # wikipedia standardizing
standard score.
# z = (xbar - p) * np.sqrt(n / (p * (1 - p)))
note = ""
Ref: https://www.statisticshowto.com/probability-and-statistics/hypothesis-testing/

z = phat - p
    -----
    sqrt( pq/n)


At 95% confidence level (alpha = 0.05):
zstar = 1.96 for two-tailed
       = 1.645 for one-tailed

if z > zstar: reject H0
""""

# from the z-score, compute the p-value
pval = 2 * (1 - stats.norm.cdf(z)) # Remember: 1-cdf gives outer zone
pval # 0.0278 # 2 means two sides of normal curve

# decision
This p-value is less than 0.05, so we reject the null hypothesis and
conclude that the coin is probably not fair.
```

After doing the test, our p-value lies outer than 5% critical region (darkblue,) and p-value (prob value) does not lie inside 95% region, so our probability is not reliable under 5% significance.



## How it works...

---

The coin tossing experiment is modeled as a sequence of  $n$  independent random variables  $x_i \in \{0, 1\}$  following the Bernoulli distribution  $B(q)$ . Each  $x_i$  represents one coin flip. After our experiment, we get actual values (samples) for these variables. A different notation is sometimes used to distinguish between the random variables (probabilistic objects) and the actual values (samples).

The following formula gives the sample mean (proportion of heads here):

$$\bar{x} = \frac{1}{n} \sum_i x_i$$

Knowing the expectancy  $\mu = q$  and variance  $\sigma^2 = q(1 - q)$  of the distribution  $B(q)$ , we compute:

$$\begin{aligned} E[\bar{x}] &= \mu = q \\ \text{var}(\bar{x}) &= \frac{\sigma^2}{n} = \frac{q(1 - q)}{n} \end{aligned}$$


The z-test is the normalized version of  $\bar{x}$  (we remove its mean, and divide by the standard deviation, thus we get a variable with mean 0 and standard deviation 1):

$$z = \frac{\bar{x} - E[\bar{x}]}{\text{std}(\bar{x})} = (\bar{x} - q) \sqrt{\frac{n}{q(1 - q)}}$$

Under the null hypothesis, what is the probability of obtaining a z-test higher (in absolute value) than some quantity  $z_0$ ? This probability is called the (two-sided) p-value. According to the central limit theorem, the z-test approximately follows a standard Gaussian distribution  $N(0, 1)$  for large  $n$ , so we get:

$$p = P[|z| > z_0] = 2P[z > z_0] \simeq 2(1 - \Phi(z_0))$$

The following diagram illustrates the z-score and the p-value:



In this formula,  $\Phi$  is the cumulative distribution function of a standard normal distribution. In SciPy, we can get it with `scipy.stats.norm.cdf`. So, given the z-test computed from the data, we compute the p-value: the probability of observing a z-test more extreme than the observed test, under the null hypothesis.

If the p-value is less than five percent (a frequently-chosen significance level, for arbitrary and historical reasons), we conclude that either:

- The null hypothesis is false, thus we conclude that the coin is unfair.
- The null hypothesis is true, and it's just bad luck if we obtained these values. We cannot make a conclusion.

We cannot disambiguate between these two options in this framework, but typically the first option is chosen. We hit the limits of frequentist statistics, although there are ways to mitigate this problem (for example, by conducting several independent studies and looking at all of their conclusions).

## One sample t-test

The One Sample t Test determines whether the sample mean is statistically different from a known or hypothesised population mean. The One Sample t Test is a parametric test. Example :- you have 50 ages and you are checking whether avg age is 30 or not.

```
from scipy.stats import ttest_1samp
import numpy as np

x = np.random.randint(20,40,size=(50))
x_mean = np.mean(x)
pop_mean = 30 # declared by researcher
tstat, pval = ttest_1samp(x, pop_mean)
alpha = 0.05

print('data = ', x)
print()
print('mean = ', x_mean)
print('p-value = ', pval)
```



```
print('alpha = ', alpha)

if pval < alpha:    # alpha value is 0.05 or 5%
    print(" we reject the null hypothesis")
else:
    print("we accept the null hypothesis")
```

## Two sampled T-test

The Independent Samples t Test or 2-sample t-test compares the means of two independent groups in order to determine whether there is statistical evidence that the associated population means are significantly different. The Independent Samples t Test is a parametric test. This test is also known as: Independent t Test.

Example : is there any association between week1 and week2

```
import numpy as np
from scipy.stats import ttest_ind # independent
np.random.seed(100)

x1 = np.random.randint(20,40,size=(50)) # randint second number is not
included
x2 = np.random.randint(20,40,size=(50)) # here, 40 is NOT included
alpha = 0.05

tstat,pval = ttest_ind(x1,x2) # independent
print("p-value",pval)
if pval <0.05:
    print("we reject null hypothesis")
else:
    print("we accept null hypothesis")

# p-value 0.9045991594475291
# we accept null hypothesis
```

## Paired sampled t-test

The paired sample t-test is also called dependent sample t-test. It's an uni-variate test that tests for a significant difference between 2 related variables. An example of this is if you where to collect the blood pressure for an individual before and after some treatment, condition, or time point.

H0 :- means difference between two sample is 0 H1 :- means difference between two sample is not 0

```
import numpy as np
from scipy import stats
np.random.seed(100)

x1 = np.random.randint(20,40,size=(50))
x2 = np.random.randint(20,40,size=(50))
```

```
alpha = 0.05
tstat,pval = stats.ttest_rel(x1, x2)
print(pval)
if pval<0.05:
    print("reject null hypothesis")
else:
    print("accept null hypothesis")
```

## One-sampled Z-test

- `statsmodels: ztest`
- `scipy.stats` has t-test but it does not have ztest, to do ztest we need to use `statsmodels.stats` and then `weightstats`

```
from scipy import stats
from statsmodels.stats import weightstats as wstats
np.random.seed(100)

x = np.random.randint(20,40,size=(50)) # weights of mice

# H0: mean(x1) - value = 0 (two-sided test)
zstat,pval = wstats.ztest(x, x2=None, value=30) # value is population mean
print(float(pval))
if pval< alpha:
    print("reject null hypothesis")
else:
    print("accept null hypothesis")

# 0.4819734873861248
# fail to reject null hypothesis
```

## Two-sampled Z-test

- `scipy.stats` has t-test but it does not have ztest, to do ztest we need to use `statsmodels.stats` and then `weightstats`

```
from scipy import stats
from statsmodels.stats import weightstats as wstats
np.random.seed(100)

x1 = np.random.randint(20,40,size=(50)) # blood sugar before
x2 = np.random.randint(20,40,size=(50)) # blood sugar after
alpha = 0.05

# value is difference of mean
zstat,pval = wstats.ztest(x1, x2=x2, value=0,alternative='two-sided')
print(f'p-value = {pval:.4f}')
if pval< alpha:
```

```
print("We Reject the Null Hypothesis.")
else:
    print("We Accept the Null Hypothesis.")
```

## F-test or ANOVA

When we have more than two groups we use ANOVA test.

## Calculate zscore (standardization)

- [scipy.stats.zscore](#)

```
import numpy as np
from scipy import stats

a = np.array([ 0.7972,  0.0767,  0.4383,  0.7866,  0.8091,
               0.1954,  0.6307,  0.6599,  0.1065,  0.0508])

a_normalized = stats.zscore(a)
print(a_normalized)

a_normalized2 = (a - a.mean()) / a.std()
print(a_normalized2) # both are same

ans = [ 1.12724554 -1.2469956 -0.05542642  1.09231569  1.16645923
        -0.8558472
        0.57858329  0.67480514 -1.14879659 -1.33234306]
```

## Binomial Test

- [scipy.stats.binomtest](#)

A car manufacturer claims that no more than 10% of their cars are unsafe. 15 cars are inspected for safety, 3 were found to be unsafe. Test the manufacturer's claim:

```
# H0: no more than 10% are unsafe (usual thing is that cars are safe)
# Ha: more than 10% are unsafe (right sided test) (alternative looks for
unsafe)

result = stats.binomtest(k=3, n=15, p=0.1, alternative='greater') # 0.184
pval = result.pvalue
print(result.proportion_estimate) # k/n
print(result.proportion_ci(confidence_level=0.95))

#The null hypothesis cannot be rejected at the 5% level of significance
# because the returned p-value is greater than the critical value of 5%.

if p < alpha:
    print("The null hypothesis can be rejected")
```

```

    print("Cars are unsafe. Status Quo is challenged by tester.")
else:
    print("The null hypothesis cannot be rejected")
    print("Cars are safe. Usual thing stands.")

```

- In a coin toss 5 times we get 4 heads. Is the coin biased?

```

import numpy as np
from scipy import stats

alpha = 0.05
result = stats.binomtest(k=4, n=5, p=0.5, alternative='two-sided')
pval = result.pvalue
print('pvalue = ', pval) # 0.375
if pval < alpha:
    print("We Reject the Null Hypothesis.")
else:
    print("We Accept the Null Hypothesis.")

# pvalue = 0.375
# We fail to reject the Null Hypothesis.

# p-value = P(4H1T) + P(4T1H) + P(5H) + P(5T)
pval2 = 5/32 + 5/32 + 1/32 + 1/32
print(pval2) # 0.375

```

## Scipy normaltest (D'Agostino and Pearson's test for normality)

`scipy` does not have `ztest` but it has some normality test. The default normal test function uses D'Agostino test.

- [scipy.stats.normaltest](#)

```

import numpy as np
from scipy import stats
np.random.seed(100)

N = 1000
a = np.random.normal(0, 1, size=N)
b = np.random.normal(2, 1, size=N)

x = np.concatenate((a, b))
k2, p = stats.normaltest(x, nan_policy='omit')

alpha = 0.001 # 1e-3 is 0.1%
print("p = {:g}".format(p))

if p < alpha: # null hypothesis: x comes from a normal distribution
    print("The null hypothesis can be rejected")

```

```
        print("Data is NOT normal.")
    else:
        print("The null hypothesis cannot be rejected")
        print("Data is normal.")

# p = p = 3.17699e-08
# The null hypothesis can be rejected

stat,p = stats.normaltest(a)
print("Reject Null Hypothesis (NOT normal)" if p < 0.05 else "Fail to
reject Null Hypothesis (normal).")
# Fail to reject Null Hypothesis (normal).
```